

```

int open(const char* path_name, int flags, ...);
ssize_t read(int fd, void* buffer, ssize_t buffer_size);
ssize_t write(int fd, void* buffer, ssize_t buffer_size);
int close(int fd);
#-----

const char* path_name = "/home/yogeshwar/src/abc.txt"
int fd = open(path_name, O_RDONLY);
assert(fd != -1);
#-----

pqr.txt does not exist in /home/yogeshwar/src
const char* path_name = "/home/yogeshwar/src/pqr.txt";
int fd = open(path_name, O_WRONLY | O_CREAT, S_IRUSR | S_IWUSR | S_IRGRP);
assert(fd != -1);
#-----

pqr.txt may or may not exist in /home/yogeshwar/src
if does not exist -> creat it
if exists -> truncate the existing content and open it

int fd = open(path_name, O_WRONLY | O_CREAT | O_TRUNC, S_IRUSR | S_IWUSR |
              S_IRGRP | S_IROTH);
assert(fd != -1)
#-----

hello.c is present in src.
You want to open it to overwrite its content
const char* path_name = "/home/yogeshwar/src/hello.c"
int fd = open(path_name, O_WRONLY);
#-----

hello.c is present in src.
You want to append data to hello.c
You want to ensure that whenever a write system call is issued, it (= write)
will make sure to seek the next read write offset to the end of the file
and then commit a write
int fd = open(path_name, O_WRONLY | O_APPEND);
assert(fd != -1)
#-----

ssize_t pread(int fd, void* buffer, size_t size, off_t offset);
ssize_t pwrite(int fd, void* buffer, size_t size, off_t offset);

OFFSET RANGE -> offset to offset + size -1 -> bmap()


off_t start_offset = 8 * 1024;
ssize_t size = 4 * 1024;
char buffer[4096];

```

buffer will contain a data from 8K to 12K-1 (assuming file is at least as long as 12K bytes)

```
off_t lseek(int fd, off_t offset, int whence);
```

offset negative -> LEFTWARDS



## LSEEK PATTERNS:

Move the next read write offset at the start of the file.

```
lseek(fd, 0, SEEK_SET);
```

Position next read write offset at the end of the file.

```
lseek(fd, 0, SEEK_END);
```

(better use O\_APPEND so that job of this lseek call is internally performed by write() avoiding the race condition)

Find out where the current offset is

```
off_t offset = lseek(fd, 0, SEEK_CUR);
```

```
void read_test(const char* path_name){
    int fd = open(path_name, O_RDONLY);
    off_t current_offset;
    char* buffer = (char*)malloc(4096);
    ssize_t read_bytes;
    while((read_bytes = read(fd, buffer, 4096)) > 0){
        current_offset = lseek(fd, 0, SEEK_CUR);
        printf("Current offset = %lu\n", current_offset);
        write(STDOUT_FILENO, buffer, read_bytes);
    }
}
```

```
void read_test(const char* path_name){
    int fd = open(path_name, O_RDONLY);
    off_t current_offset;
    char* buffer = (char*)malloc(4096);
    ssize_t read_bytes;
    off_t off_read = 0;
    while((read_bytes = pread(fd, buffer, 4096, off_read)) > 0){
        current_offset = lseek(fd, 0, SEEK_CUR);
        printf("Current offset = %lu\n", current_offset);
        off_read += read_bytes;
        write(STDOUT_FILENO, buffer, read_bytes);
    }
}
```

/\* Assume that




1) LOGICAL BLOCK SIZE == 4K

2) File is being processed only by this application (no race condition)





3) Initial size of file = 8K

\*/



```
void test(const char* path_name){  
    int fd = open(path_name, O_RDWR);  
    off_t offset = lseek(fd, 0, SEEK_END);  
    printf("END OFFSET == SIZE OF FILE = %lu\n", offset);  
    offset = lseek(fd, 4096, SEEK_END);  
    write(fd, "#", 1); /* Write # character to file */  
    close(fd);  
}
```




inode, 13 member array



$\xrightarrow{\text{lseek}(fd, 0, \text{SEEK\_END})}$






$\xrightarrow{\text{lseek}(fd, 4096, \text{SEEK\_END})}$



12K = #

#







`readv() / writev()`


SCATTER / GATHER - I/O

`ssize_t readv(int fd, const struct iovec* iov, int iovcnt);`

`ssize_t writev(int fd, const struct iovec* iov, int iovcnt);`



{  
 void\* buffer  
 size\_t size  
 }



// BIG FILE

```

void test(const char* path_name){
    static struct iovec iov_arr[4];
    int i;

    for(i = 0; i < 4; ++i){
        iov_arr[i].iov_base = (char*)malloc(4096);
        iov_arr[i].iov_len = 4096;
    }

    int fd = open(path_name, O_RDONLY);
    assert(fd != -1);

    int read_bytes = readv(fd, iov_arr, 4);
    // 16 KB have been read
    // first 4 KB -> iov_arr[0].iov_base
    // second 4 KB -> iov_arr[1].iov_base
    // third 4 KB -> iov_arr[2].iov_base
    // fourth 4 KB -> iov_arr[3].iov_base
    for(i = 0; i < 4; ++i)
        printf("ARRAY(%d):%s\n", i, (char*)iov_arr[i].iov_base);
    for(i = 0; i < 4; ++i)
        free(iov_arr[i].iov_base);
}
  
```

`open(), read(), write(), lseek(), close(), pread(), pwrite(),  
readv(), writev(), preadv(), pwritev()`

`#-----`

`fcntl(), ioctl(),  
dup(), dup2(), dup3()`

`#-----`

`fork(), execve(), wait(), _exit()`

`#-----`

**Basic signal**